


Python Data Structures implementation

list, dict: how does CPython actually implement them?

Flavien Raynaud


05 February 2017

FOSDEM 2017



Yelp's Mission

Connecting people with great local businesses.



list & dict

► To get started

```
>>> repeat('l.append(42)', 'l = [ ]')  
2.1469707062351516e-07  
>>> repeat('l.append(42)', 'l = [1]')  
2.2675518121104688e-07  
>>> repeat('l.append(42)', 'l = [i for i in range(1000)]')  
2.475244084052974e-07
```

```
>>> repeat('d["year"] = 2017', 'd = {}')  
1.5317109500756487e-07  
>>> repeat('d["year"] = 2017', 'd = {"one": 1}')  
1.5375611219496933e-07  
>>> repeat('d["year"] = 2017', 'd = {str(i): i for i in range(1000)}')  
2.46819700623746e-07
```

list & dict

- ▶ Focus on CPython 3.6
- ▶ A lot of hidden (and really cool) ideas
- ▶ A lot of lines of code (and comments)
 - ▶ ~3500 for lists
 - ▶ ~4500 for dicts
- ▶ (Almost) everyone knows (at least roughly) how they work

- ▶ A sequence of values (read: objects), 0-indexed
- ▶ $\mathcal{O}(1)$ amortized insert, $\mathcal{O}(1)$ random access, $\mathcal{O}(n)$ deletion

- ▶ A sequence of values (read: objects), 0-indexed
- ▶ $\mathcal{O}(1)$ amortized insert, $\mathcal{O}(1)$ random access, $\mathcal{O}(n)$ deletion
- ▶ Vector
 - ▶ Over-allocated array
 - ▶ Invariant: $0 \leq \text{len}(\text{list}) \leq \text{capacity}$

creating a new list

```
list()    # please avoid :-)  
[]  
[0, 1, 2, 3, 4]  
list((0, 1, 2, 3, 4))  
[i for i in range(5)]
```

- ▶ `[]`, `list()` → size = 0, capacity = 0
- ▶ `[0, 1, 2, 3, 4]` → size = 5, capacity = 5

appending to a list

```
categories = []  
categories.append('food')
```


appending to a list

```
categories = []  
categories.append('food')
```

What is really happening?

- ▶ `resize(size+1)`
- ▶ set last value to `'food'`

resize

```
# resize the vector if necessary
resize(new_size):
    if capacity/2 <= new_size <= capacity:
        return

    capacity = (new_size / 8) + new_size
    capacity += (3 if new_size < 9 else 6)
    # realloc
```

resize

```
# resize the vector if necessary
resize(new_size):
    if capacity/2 <= new_size <= capacity:
        return

    capacity = (new_size / 8) + new_size
    capacity += (3 if new_size < 9 else 6)
    # realloc
```

- ▶ 0, 4, 8, 16, 25, 35, 46, 58, 72, 88, ...
- ▶ Growth rate: ~12.5%

append takeaway

```
>>> categories = [  
    'food', 'tacos', 'bar', 'dentist', 'scuba diving'  
]  
>>> size_capacity(categories)  
SizeCapacity(size=5, capacity=5)
```

append takeaway

```
>>> categories = [  
    'food', 'tacos', 'bar', 'dentist', 'scuba diving'  
]  
>>> size_capacity(categories)  
SizeCapacity(size=5, capacity=5)
```

```
>>> categories = []  
>>> categories.append('food')  
>>> categories.append('tacos')  
>>> categories.append('bar')  
>>> categories.append('dentist')  
>>> categories.append('scuba diving')  
>>> size_capacity(categories)  
SizeCapacity(size=5, capacity=8)
```

removing from a list

```
categories = ['food', 'tacos', 'bar', 'dentist']  
categories.pop()  
categories.pop(i)
```

- ▶ `i == size - 1`
 - ▶ `resize(size - 1)`
- ▶ `i < size - 1`
 - ▶ `categories[i:] = categories[i+1:]`
 - ▶ `memmove, resize(size - 1)`

removing from a list - $i < \text{size} - 1$

```
categories = ['food', 'tacos', 'bar', 'dentist']  
categories.pop(1)  # no more tacos :-(
```

food	tacos	bar	dentist
------	-------	-----	---------

size = 4, capacity = 4

removing from a list - $i < \text{size} - 1$

```
categories = ['food', 'tacos', 'bar', 'dentist']  
categories.pop(1)  # no more tacos :-(
```

food	bar	dentist	dentist
------	-----	---------	---------

size = 4, capacity = 4

removing from a list - $i < \text{size} - 1$

```
categories = ['food', 'tacos', 'bar', 'dentist']  
categories.pop(1)  # no more tacos :-)
```

food	bar	dentist	dentist
------	-----	---------	---------

size = 4, capacity = 4

removing from a list - $i < \text{size} - 1$

```
categories = ['food', 'tacos', 'bar', 'dentist']  
categories.pop(1)  # no more tacos :-(
```

food	bar	dentist	
------	-----	---------	--

size = 3, capacity = 4

list misc.

- ▶ `list` as a queue (`.append()`, `.pop(0)`) is bad → `deque`

list misc.

- ▶ list as a queue (`.append()`, `.pop(0)`) is bad → deque
- ▶ slicing is really powerful!

```
ints = [0, 1, 2, 3, 4]
ints[1:4] = [42] -> [0, 42, 4]
ints[1:1] = [42, 43] -> [0, 42, 43, 1, 2, 3, 4]
```

list misc.

- ▶ list as a queue (`.append()`, `.pop(0)`) is bad → deque
- ▶ slicing is really powerful!

```
ints = [0, 1, 2, 3, 4]
ints[1:4] = [42] -> [0, 42, 4]
ints[1:1] = [42, 43] -> [0, 42, 43, 1, 2, 3, 4]
```

- ▶ reference reuse scheme

```
>>> a, b, c = [0, 1], [2, 3], [4, 5]
>>> id(a), id(b), id(c)
(140512822066120, 140512822065864, 140512822065928)
>>> del b
>>> d = [6, 7]
>>> id(d)
140512822065864
```

dict

- ▶ `dict` = dictionary
- ▶ Store (*key*, *value*) pairs

creating a new dict

```
dict()    # please avoid :-)
{}
{str(i): i for i in range(5)}
dict([('1', 1), ('2', 2)])
{
    'name': 'flavr',
    'nationality': 'french',
    'language': 'python',
    'age': 42,
}
```

dict usecases

- ▶ kwargs
 - ▶ ~1 write, ~1 read, small length

dict usecases

- ▶ `kwargs`
 - ▶ ~1 write, ~1 read, small length
- ▶ `class methods (MyClass.__dict__)`
 - ▶ ~1 write, many reads, any length but all share 8-16 elements

dict usecases

- ▶ `kwargs`
 - ▶ ~1 write, ~1 read, small length
- ▶ `class methods (MyClass.__dict__)`
 - ▶ ~1 write, many reads, any length but all share 8-16 elements
- ▶ `attributes, global vars (obj.__dict__, globals())`
 - ▶ many writes, many reads, any length but often < 10

dict usecases

- ▶ `kwargs`
 - ▶ ~1 write, ~1 read, small length
- ▶ `class methods (MyClass.__dict__)`
 - ▶ ~1 write, many reads, any length but all share 8-16 elements
- ▶ `attributes, global vars (obj.__dict__, globals())`
 - ▶ many writes, many reads, any length but often < 10
- ▶ `builtins (__builtins__.__dict__)`
 - ▶ ~0 writes, many reads, length ~150

dict usecases

- ▶ `kwargs`
 - ▶ ~1 write, ~1 read, small length
- ▶ `class methods (MyClass.__dict__)`
 - ▶ ~1 write, many reads, any length but all share 8-16 elements
- ▶ `attributes, global vars (obj.__dict__, globals())`
 - ▶ many writes, many reads, any length but often < 10
- ▶ `builtins (__builtins__.__dict__)`
 - ▶ ~0 writes, many reads, length ~150
- ▶ `uniquification (remove duplicates, counters)`
 - ▶ many writes, ~1 read, any length

dict usecases

- ▶ `kwargs`
 - ▶ ~1 write, ~1 read, small length
- ▶ `class methods (MyClass.__dict__)`
 - ▶ ~1 write, many reads, any length but all share 8-16 elements
- ▶ `attributes, global vars (obj.__dict__, globals())`
 - ▶ many writes, many reads, any length but often < 10
- ▶ `builtins (__builtins__.__dict__)`
 - ▶ ~0 writes, many reads, length ~150
- ▶ `uniquification (remove duplicates, counters)`
 - ▶ many writes, ~1 read, any length
- ▶ `other use`
 - ▶ any writes, any reads, any length, any deletions

dict history

- ▶ many implementation changes (2.1, 3.3, 3.6)
- ▶ dict in CPython 3.6
 - ▶ inspired from Pypy
 - ▶ ordered (`.keys()`, `.values()`, `.items()`)
 - ▶ memory-efficient (re-use keys when possible)
 - ▶ PEP412 - Key-Sharing dictionary
 - ▶ Split table, **Combined table**

dict

- ▶ $\mathcal{O}(1)$ average insert, $\mathcal{O}(1)$ average lookup, $\mathcal{O}(1)$ average deletion
- ▶ Fast access? Arrays
- ▶ dict key \leftrightarrow array index

dict

- ▶ $\mathcal{O}(1)$ average insert, $\mathcal{O}(1)$ average lookup, $\mathcal{O}(1)$ average deletion
- ▶ Fast access? Arrays
- ▶ dict key \leftrightarrow array index
- ▶ Hashing

hashing

Hash function: function used to map data from arbitrary size to data of (almost always) fixed size.

- ▶ CPython: {32,64}-bit integers

hashing

Hash function: function used to map data from arbitrary size to data of (almost always) fixed size.

- CPython: {32,64}-bit integers

[illegible]

hashing

Hash function: function used to map data from arbitrary size to data of (almost always) fixed size.

- CPython: {32,64}-bit integers

[illegible][illegible]

hashing

- Similar values often have dissimilar hashes

```
>>> bits(hash("hello"))  
'0011010000111000011101000100111110100011110011111010100110100000'  
>>> bits(hash("hallo"))  
'1110100100100011001111111101011101100111100001100011110011101011'
```

hashing

- ▶ Similar values often have dissimilar hashes

```
>>> bits(hash("hello"))
'0011010000111000011101000100111110100011110011111010100110100000'
>>> bits(hash("hallo"))
'1110100100100011001111111101011101100111100001100011110011101011'
```

- ▶ Same value \Rightarrow same hash

```
>>> bits(hash("hello"))
'0011010000111000011101000100111110100011110011111010100110100000'
>>> bits(hash("hello"))
'0011010000111000011101000100111110100011110011111010100110100000'
>>> bits(hash("hello"))
'0011010000111000011101000100111110100011110011111010100110100000'
```

dict

- ▶ dict key \rightarrow key hash \rightarrow array index

dict

- ▶ dict key \rightarrow key hash \rightarrow array index
- ▶ Can we actually represent a dict using arrays?

dict

- ▶ dict key \rightarrow key hash \rightarrow array index
- ▶ Can we actually represent a dict using arrays?
 - ▶ Yes.
 - ▶ dict = 2 arrays (indices, entries)

dict as arrays

```
# categories: dict(key=name, value=#businesses)  
categories = {}
```

dict as arrays

```
# categories: dict(key=name, value=#businesses)
categories = {}
```

		entry index
0	000	
1	001	
2	010	
3	011	
4	100	
5	101	
6	110	
7	111	

	hash	key	value

► initial size = 8

dict & hash

- ▶ index of key x
 - ▶ last bits of hash(x)

```
>>> categories['food'] = 4000
>>> bits(hash('food'))
'0111100011101101110010011011100110110110110000001011001001001000'
>>> bits(hash('food'))[-3:]
'000'
```


inserting in a dict

```
>>> categories['tacos'] = 31
>>> bits(hash('tacos'))[-3:]
'001'
>>> categories['bar'] = 127
>>> bits(hash('bar'))[-3:]
'101'
```

inserting in a dict

```
>>> categories['tacos'] = 31
>>> bits(hash('tacos'))[-3:]
'001'
>>> categories['bar'] = 127
>>> bits(hash('bar'))[-3:]
'101'
```

		entry index
0	000	0
1	001	1
2	010	
3	011	
4	100	
5	101	
6	110	
7	111	

	hash	key	value
0	01...000	'food'	4000
1	10...001	'tacos'	31

inserting in a dict

```
>>> categories['tacos'] = 31
>>> bits(hash('tacos'))[-3:]
'001'
>>> categories['bar'] = 127
>>> bits(hash('bar'))[-3:]
'101'
```

		entry index
0	000	0
1	001	1
2	010	
3	011	
4	100	
5	101	2
6	110	
7	111	

	hash	key	value
0	01...000	'food'	4000
1	10...001	'tacos'	31
2	00...101	'bar'	127

inserting in a dict

```
>>> categories['dentist'] = 17
>>> bits(hash('dentist'))[-3:]
'001'
```


inserting in a dict

```
>>> categories['dentist'] = 17  
>>> bits(hash('dentist'))[-3:]  
'001'
```

		entry index
0	000	0
1	001	1
2	010	
3	011	
4	100	
5	101	2
6	110	
7	111	

	hash	key	value
0	01...000	'food'	4000
1	10...001	'tacos'	31
2	00...101	'bar'	127

collision resolution

- ▶ Hash collision resolution: Open Addressing
 - ▶ $index = (5 * index + 1) \% size$

collision resolution

- ▶ Hash collision resolution: Open Addressing
 - ▶ $index = (5 * index + 1) \% size$
 - ▶ traverses each integer in $\{0, \dots, size - 1\}$
 - ▶ (actually a bit more sophisticated)

inserting in a dict

- ▶ $index = 001_2 = 1$

inserting in a dict

- ▶ $index = 001_2 = 1$
- ▶ $index = (5 \times 1 + 1) \% 8 = 6 = 110_2$

inserting in a dict

- ▶ $index = 001_2 = 1$
- ▶ $index = (5 \times 1 + 1) \% 8 = 6 = 110_2$

		entry index
0	000	0
1	001	1
2	010	
3	011	
4	100	
5	101	2
6	110	3
7	111	

	hash	key	value
0	01...000	'food'	4000
1	10...001	'tacos'	31
2	00...101	'bar'	127
3	11...001	'dentist'	17

lookup in a dict

```
>>> categories['food']  
4000  
>>> bits(hash('food'))[-3:]  
'000'
```

lookup in a dict

```
>>> categories['food']  
4000  
>>> bits(hash('food'))[-3:]  
'000'
```

		entry index
0	000	0
1	001	1
2	010	
3	011	
4	100	
5	101	2
6	110	3
7	111	

	hash	key	value
0	01...000	'food'	4000
1	10...001	'tacos'	31
2	00...101	'bar'	127
3	11...001	'dentist'	17

lookup in a dict

```
>>> categories['dentist']  
17  
>>> bits(hash('dentist'))[-3:]  
'001'
```

lookup in a dict

```
>>> categories['dentist']  
17  
>>> bits(hash('dentist'))[-3:]  
'001'
```

		entry index
0	000	0
1	001	1
2	010	
3	011	
4	100	
5	101	2
6	110	3
7	111	

	hash	key	value
0	01...000	'food'	4000
1	10...001	'tacos'	31
2	00...101	'bar'	127
3	11...001	'dentist'	17

lookup in a dict

```
>>> categories['dentist']  
17  
>>> bits(hash('dentist'))[-3:]  
'001'
```

		entry index
0	000	0
1	001	1
2	010	
3	011	
4	100	
5	101	2
6	110	3
7	111	

	hash	key	value
0	01...000	'food'	4000
1	10...001	'tacos'	31
2	00...101	'bar'	127
3	11...001	'dentist'	17

lookup in a dict

```
>>> categories['music']
Traceback (most recent call last):
...
KeyError: 'music'
>>> bits(hash('music'))[-3:]
'000'
```

lookup in a dict

```
>>> categories['music']
Traceback (most recent call last):
...
KeyError: 'music'
>>> bits(hash('music'))[-3:]
'000'
```

		entry index
0	000	0
1	001	1
2	010	
3	011	
4	100	
5	101	2
6	110	3
7	111	

	hash	key	value
0	01...000	'food'	4000
1	10...001	'tacos'	31
2	00...101	'bar'	127
3	11...001	'dentist'	17

lookup in a dict

```
>>> categories['music']
Traceback (most recent call last):
...
KeyError: 'music'
>>> bits(hash('music'))[-3:]
'000'
```

`next_index('000') = '001'`

		entry index
0	000	0
1	001	1
2	010	
3	011	
4	100	
5	101	2
6	110	3
7	111	

	hash	key	value
0	01...000	'food'	4000
1	10...001	'tacos'	31
2	00...101	'bar'	127
3	11...001	'dentist'	17

lookup in a dict

```
>>> categories['music']
Traceback (most recent call last):
...
KeyError: 'music'
>>> bits(hash('music'))[-3:]
'000'
```

`next_index('001') = '110'`

		entry index
0	000	0
1	001	1
2	010	
3	011	
4	100	
5	101	2
6	110	3
7	111	

	hash	key	value
0	01...000	'food'	4000
1	10...001	'tacos'	31
2	00...101	'bar'	127
3	11...001	'dentist'	17

lookup in a dict

```
>>> categories['music']
Traceback (most recent call last):
...
KeyError: 'music'
>>> bits(hash('music'))[-3:]
'000'
```

`next_index('110') = '111'`

		entry index
0	000	0
1	001	1
2	010	
3	011	
4	100	
5	101	2
6	110	3
7	111	

	hash	key	value
0	01...000	'food'	4000
1	10...001	'tacos'	31
2	00...101	'bar'	127
3	11...001	'dentist'	17

deleting from a dict

```
>>> del categories['tacos']  
>>> bits(hash('tacos'))[-3:]  
'001'
```

deleting from a dict

```
>>> del categories['tacos']  
>>> bits(hash('tacos'))[-3:]  
'001'
```

		entry index
0	000	0
1	001	
2	010	
3	011	
4	100	2
5	101	
6	110	
7	111	

	hash	key	value
0	01...000	'food'	4000
1			
2	00...101	'bar'	127
3	11...001	'dentist'	17

deleting from a dict

```
>>> del categories['tacos']  
>>> bits(hash('tacos'))[-3:]  
'001'
```

		entry index
0	000	0
1	001	
2	010	
3	011	
4	100	2
5	101	
6	110	
7	111	

	hash	key	value
0	01...000	'food'	4000
1			
2	00...101	'bar'	127
3	11...001	'dentist'	17

- 'dentist' is not accessible anymore!

deleting from a dict

```
>>> del categories['tacos']  
>>> bits(hash('tacos'))[-3:]  
'001'
```

		entry index
0	000	0
1	001	1
2	010	
3	011	
4	100	
5	101	2
6	110	3
7	111	

	hash	key	value
0	01...000	'food'	4000
1		<dummy>	
2	00...101	'bar'	127
3	11...001	'dentist'	17

deleting from a dict

```
>>> del categories['tacos']  
>>> bits(hash('tacos'))[-3:]  
'001'
```

		entry index
0	000	0
1	001	1
2	010	
3	011	
4	100	
5	101	2
6	110	3
7	111	

	hash	key	value
0	01...000	'food'	4000
1		<dummy>	
2	00...101	'bar'	127
3	11...001	'dentist'	17

- 'dentist' is still accessible!

caveats

- ▶ 8 slots is rarely enough
- ▶ full `indices` array → slower lookups
- ▶ `< dummy >` keys → even slower

resizing a dict

- ▶ invariant: at least one empty slot
- ▶ $\text{usable} = \frac{2}{3} \text{ size}$ (= 5 initially)

resizing a dict

- ▶ invariant: at least one empty slot
- ▶ $\text{usable} = \frac{2}{3} \text{ size}$ (= 5 initially)

		entry index
0	000	0
1	001	1
2	010	
3	011	
4	100	
5	101	2
6	110	3
7	111	

	hash	key	value
0	01...000	'food'	4000
1		<dummy>	
2	00...101	'bar'	127
3	11...001	'dentist'	17

- ▶ `len = 3, size = 8, usable = 1`

resizing a dict

```
>>> categories['dinner'] = 1024  
>>> del categories['dentist']
```

resizing a dict

```
>>> categories['dinner'] = 1024  
>>> del categories['dentist']
```

		entry index
0	000	0
1	001	1
2	010	
3	011	4
4	100	
5	101	2
6	110	3
7	111	

	hash	key	value
0	01...000	'food'	4000
1		<dummy>	
2	00...101	'bar'	127
3		<dummy>	
4	00...011	'dinner'	1024

► len = 3, size = 8, usable = 0

resizing a dict

```
>>> categories['vegan'] = 1024
```

resizing a dict

```
>>> categories['vegan'] = 1024
```

► `resize(2 * len + size / 2)`

```
resize(min_size):  
    # to truncate hashes  
    new_size = NEXT_POWER_OF_TWO(min_size)  
    create_new_dict(new_size)  
    for (hash, key, value) in entries:  
        insert_new(hash, key, value)  
    delete_old_dict()
```

► `len = 3, size = 8` → `min_size = 10, new_size = 16`

resizing a dict

```
>>> categories['vegan'] = 1024
```

► `resize(2 * len + size / 2)`

```
resize(min_size):  
    # to truncate hashes  
    new_size = NEXT_POWER_OF_TWO(min_size)  
    create_new_dict(new_size)  
    for (hash, key, value) in entries:  
        insert_new(hash, key, value)  
    delete_old_dict()
```

- `len = 3, size = 8` → `min_size = 10, new_size = 16`
- larger arrays → more items can fit
- larger arrays → more free slots → faster lookups
- no more *< dummy >* keys

- ▶ reference reuse scheme
- ▶ split table
 - ▶ 3 arrays (indices, entries, values)
 - ▶ share (indices, entries), own values

References/Resources:

- ▶ `github.com/flavray/fosdem-python-data-structures`
- ▶ CPython 3.6
- ▶ PEP412 - Key-Sharing Dictionary
- ▶ Faster, more memory efficient and more ordered dictionaries on PyPy
- ▶ The Mighty Dictionary (2010) - Brandon Craig Rhodes

special cases

```
>>> ints = [0, 1, 2, 3, 4]
>>> size_capacity(ints)
SizeCapacity(size=5, capacity=5)
```

```
>>> ints = [i for i in range(5)]
# Comprehensions act as for-loops, .append
>>> size_capacity(ints)
SizeCapacity(size=5, capacity=8)
```


ordering

```
>>> categories.keys(), categories.values()  
(['food', 'bar', 'dinner'], [4000, 127, 1024])
```

		entry index
0	000	0
1	001	1
2	010	
3	011	4
4	100	
5	101	2
6	110	3
7	111	

	hash	key	value
0	01...000	'food'	4000
1		<dummy>	
2	00...101	'bar'	127
3		<dummy>	
4	00...011	'dinner'	1024

ordering

```
>>> categories.keys(), categories.values()
(['food', 'bar', 'dinner'], [4000, 127, 1024])
```

		entry index
0	000	0
1	001	1
2	010	
3	011	4
4	100	
5	101	2
6	110	3
7	111	

	hash	key	value
0	01...000	'food'	4000
1		<dummy>	
2	00...101	'bar'	127
3		<dummy>	
4	00...011	'dinner'	1024

\o/